

# Constraints

Foreign Keys

Local and Global Constraints

Triggers

# Constraints and Triggers

- ◆ A *constraint* is a relationship among data elements that the DBMS is required to enforce.
  - ◆ Example: key constraints.
- ◆ *Triggers* are only executed when a specified condition occurs, e.g., insertion of a tuple.
  - ◆ Easier to implement than complex constraints.

# Kinds of Constraints

- ◆ Keys.
- ◆ Foreign-key, or referential-integrity.
- ◆ Value-based constraints.
  - ◆ Constrain values of a particular attribute.
- ◆ Tuple-based constraints.
  - ◆ Relationship among components.
- ◆ Assertions: any SQL boolean expression.

# Foreign Keys

- ◆ Consider Relation *Sells(bar, beer, price)*.
- ◆ We might expect that a beer value is a real beer --- something appearing in *Beers.name* .
- ◆ A constraint that requires a beer in *Sells* to be a beer in *Beers* is called a *foreign-key* constraint.

# Expressing Foreign Keys

- ◆ Use the keyword REFERENCES, either:
  1. Within the declaration of an attribute (only for one-attribute keys).
  2. As an element of the schema:  

```
FOREIGN KEY ( <list of attributes> )  
REFERENCES <relation> ( <attributes> )
```
- ◆ Referenced attributes must be declared PRIMARY KEY or UNIQUE.

# Example: With Attribute

```
CREATE TABLE Beers (  
    name      CHAR(20) PRIMARY KEY,  
    manf      CHAR(20) );  
  
CREATE TABLE Sells (  
    bar       CHAR(20),  
    beer      CHAR(20) REFERENCES Beers(name),  
    price     REAL );
```

# Example: As Element

```
CREATE TABLE Beers (  
    name      CHAR(20) PRIMARY KEY,  
    manf      CHAR(20) );  
  
CREATE TABLE Sells (  
    bar       CHAR(20),  
    beer      CHAR(20),  
    price     REAL,  
    FOREIGN KEY(beer) REFERENCES  
        Beers(name) );
```

# Enforcing Foreign-Key Constraints

- ◆ If there is a foreign-key constraint from attributes of relation  $R$  to a key of relation  $S$ , two violations are possible:
  1. An insert or update to  $R$  introduces values not found in  $S$ .
  2. A deletion or update to  $S$  causes some tuples of  $R$  to “dangle.”



# Actions Taken --- (1)

- ◆ Suppose  $R = \text{Sells}$ ,  $S = \text{Beers}$ .
- ◆ An insert or update to Sells that introduces a nonexistent beer must be rejected.
- ◆ A deletion or update to Beers that removes a beer value found in some tuples of Sells can be handled in three ways (next slide).

# Actions Taken --- (2)

1. *Default* : Reject the modification.
2. *Cascade* : Make the same changes in Sells.
  - ◆ Deleted beer: delete Sells tuple.
  - ◆ Updated beer: change value in Sells.
3. *Set NULL* : Change the beer to NULL.

# Example: Cascade

- ◆ Delete the Bud tuple from Beers:
  - ◆ Then delete all tuples from Sells that have beer = 'Bud'.
- ◆ Update the Bud tuple by changing 'Bud' to 'Budweiser':
  - ◆ Then change all Sells tuples with beer = 'Bud' so that beer = 'Budweiser'.

# Example: Set NULL

- ◆ Delete the Bud tuple from Beers:
  - ◆ Change all tuples of Sells that have beer = 'Bud' to have beer = NULL.
- ◆ Update the Bud tuple by changing 'Bud' to 'Budweiser':
  - ◆ Same change.

# Choosing a Policy

- ◆ When we declare a foreign key, we may choose policies SET NULL or CASCADE independently for deletions and updates.
- ◆ Follow the foreign-key declaration by:  
ON [UPDATE, DELETE][SET NULL CASCADE]
- ◆ Two such clauses may be used.
- ◆ Otherwise, the default (reject) is used.

# Example

```
CREATE TABLE Sells (  
    bar    CHAR(20),  
    beer   CHAR(20),  
    price  REAL,  
    FOREIGN KEY(beer)  
        REFERENCES Beers(name)  
        ON DELETE SET NULL  
        ON UPDATE CASCADE  
);
```

# Attribute-Based Checks

- ◆ Constraints on the value of a particular attribute.
- ◆ Add: CHECK( <condition> ) to the declaration for the attribute.
- ◆ The condition may use the name of the attribute, but any other relation or attribute name must be in a subquery.

# Example

```
CREATE TABLE Sells (  
    bar      CHAR(20) ,  
    beer     CHAR(20)    CHECK ( beer IN  
        (SELECT name FROM Beers) ) ,  
    price    REAL CHECK ( price <= 5.00 )  
);
```



# Timing of Checks

- ◆ Attribute-based checks performed only when a value for that attribute is inserted or updated.
  - ◆ **Example:** `CHECK (price <= 5.00)` checks every new price and rejects the modification (for that tuple) if the price is more than \$5.
  - ◆ **Example:** `CHECK (beer IN (SELECT name FROM Beers))` not checked if a beer is deleted from Beers (unlike foreign-keys).

# Tuple-Based Checks

- ◆ CHECK ( <condition> ) may be added as a relation-schema element.
- ◆ The condition may refer to any attribute of the relation.
  - ◆ But any other attributes or relations require a subquery.
- ◆ Checked on insert or update only.

# Example: Tuple-Based Check

- ◆ Only Joe's Bar can sell beer for more than \$5:

```
CREATE TABLE Sells (  
    bar          CHAR(20),  
    beer         CHAR(20),  
    price        REAL,  
    CHECK (bar = 'Joe's Bar' OR  
           price <= 5.00)  
);
```

# Assertions

- ◆ These are database-schema elements, like relations or views.

- ◆ Defined by:

```
CREATE ASSERTION <name>  
CHECK ( <condition> );
```

- ◆ Condition may refer to any relation or attribute in the database schema.

# Example: Assertion

- ◆ In `Sells(bar, beer, price)`, no bar may charge an average of more than \$5.


```
CREATE ASSERTION NoRipoffBars CHECK (  
  NOT EXISTS (  

```

```
    SELECT bar FROM Sells  
    GROUP BY bar  
    HAVING 5.00 < AVG(price)
```

```
  ));
```

Bars with an  
average price  
above \$5



# Example: Assertion

- ◆ In `Drinkers(name, addr, phone)` and `Bars(name, addr, license)`, there cannot be more bars than drinkers.

```
CREATE ASSERTION FewBar CHECK (  
    (SELECT COUNT(*) FROM Bars) <=  
    (SELECT COUNT(*) FROM Drinkers)  
);
```

# Timing of Assertion Checks

- ◆ In principle, we must check every assertion after every modification to any relation of the database.
- ◆ A clever system can observe that only certain changes could cause a given assertion to be violated.
  - ◆ **Example:** No change to Beers can affect FewBar. Neither can an insertion to Drinkers.

# Triggers: Motivation

- ◆ Assertions are powerful, but the DBMS often can't tell when they need to be checked.
- ◆ Attribute- and tuple-based checks are checked at known times, but are not powerful.
- ◆ Triggers let the user decide when to check for a powerful condition.



# Event-Condition-Action Rules

- ◆ Another name for “trigger” is *ECA rule*, or *event-condition-action* rule.
- ◆ *Event*: typically a type of database modification, e.g., “insert on Sells.”
- ◆ *Condition*: Any SQL boolean-valued expression.
- ◆ *Action*: Any SQL statements.

# Preliminary Example: A Trigger

- ◆ Instead of using a foreign-key constraint and rejecting insertions into `Sells(bar, beer, price)` with unknown beers, a trigger can add that beer to `Beers`, with a `NULL` manufacturer.

# Example: Trigger Definition

```
CREATE TRIGGER BeerTrig
```

```
AFTER INSERT ON Sells
```

The event

```
REFERENCING NEW ROW AS NewTuple  
FOR EACH ROW
```

```
WHEN (NewTuple.beer NOT IN  
      (SELECT name FROM Beers))
```

The condition

```
INSERT INTO Beers(name)  
VALUES(NewTuple.beer);
```

The action

# Options: CREATE TRIGGER

- ◆ CREATE TRIGGER <name>

- ◆ Option:

CREATE OR REPLACE TRIGGER <name>

- ◆ Useful if there is a trigger with that name and you want to modify the trigger.

# Options: The Event

- ◆ AFTER can be BEFORE.
  - ◆ Also, INSTEAD OF, if the relation is a view.
    - A great way to execute view modifications: have triggers translate them to appropriate modifications on the base tables.
- ◆ INSERT can be DELETE or UPDATE.
  - ◆ And UPDATE can be UPDATE . . . ON a particular attribute.

# Options: FOR EACH ROW

- ◆ Triggers are either “row-level” or “statement-level.”
- ◆ FOR EACH ROW indicates row-level; its absence indicates statement-level.
- ◆ *Row level triggers* : execute once for each modified tuple.
- ◆ *Statement-level triggers* : execute once for an SQL statement, regardless of how many tuples are modified.

# Options: REFERENCING

- ◆ INSERT statements imply a new tuple (for row-level) or new table (for statement-level).
  - ◆ The “table” is the set of inserted tuples.
- ◆ DELETE implies an old tuple or table.
- ◆ UPDATE implies both.
- ◆ Refer to these by  
[NEW OLD][TUPLE TABLE] AS <name>

# Options: The Condition

- ◆ Any boolean-valued condition is appropriate.
- ◆ It is evaluated before or after the triggering event, depending on whether BEFORE or AFTER is used in the event.
- ◆ Access the new/old tuple or set of tuples through the names declared in the REFERENCING clause.



# Options: The Action

- ◆ There can be more than one SQL statement in the action.
  - ◆ Surround by BEGIN . . . END if there is more than one.
- ◆ But queries make no sense in an action, so we are really limited to modifications.

# Another Example

- ◆ Using `Sells(bar, beer, price)` and a unary relation `RipoffBars(bar)` created for the purpose, maintain a list of bars that raise the price of any beer by more than \$1.

# The Trigger

```
CREATE TRIGGER PriceTrig
```

```
AFTER UPDATE OF price ON Sells
```

The event –  
only changes  
to prices

```
REFERENCING
```

```
OLD ROW AS ooo
```

```
NEW ROW AS nnn
```

Updates let us  
talk about old  
and new tuples

```
FOR EACH ROW
```

We need to consider  
each price change

Condition:  
a raise in  
price > \$1

```
WHEN(nnn.price > ooo.price + 1.00)
```

```
INSERT INTO RipoffBars  
VALUES(nnn.bar);
```

When the price change  
is great enough, add  
the bar to RipoffBars

# Triggers on Views

- ◆ Generally, it is impossible to modify a view, because it doesn't exist.
- ◆ But an INSTEAD OF trigger lets us interpret view modifications in a way that makes sense.
- ◆ Example: We'll design a view Synergy that has (drinker, beer, bar) triples such that the bar serves the beer, the drinker frequents the bar and likes the beer.

# Example: The View

CREATE VIEW Synergy AS

SELECT Likes.drinker, Likes.beer, Sells.bar

Pick one copy of  
each attribute

FROM Likes, Sells, Frequents

WHERE Likes.drinker = Frequents.drinker

AND Likes.beer = Sells.beer

AND Sells.bar = Frequents.bar;

Natural join of Likes,  
Sells, and Frequents

# Interpreting a View Insertion

- ◆ We cannot insert into Synergy --- it is a view.
- ◆ But we can use an INSTEAD OF trigger to turn a (drinker, beer, bar) triple into three insertions of projected pairs, one for each of Likes, Sells, and Frequents.
  - ◆ The Sells.price will have to be NULL.

# The Trigger

```
CREATE TRIGGER ViewTrig
  INSTEAD OF INSERT ON Synergy
  REFERENCING NEW ROW AS n
  FOR EACH ROW
  BEGIN
    INSERT INTO LIKES VALUES(n.drinker, n.beer);
    INSERT INTO SELLS(bar, beer) VALUES(n.bar, n.beer);
    INSERT INTO FREQUENTS VALUES(n.drinker, n.bar);
  END;
```