

# Real SQL Programming

Embedded SQL

Call-Level Interface

Java Database Connectivity

# SQL in Real Programs

- ◆ We have seen only how SQL is used at the generic query interface --- an environment where we sit at a terminal and ask queries of a database.
- ◆ Reality is almost always different.
  - ◆ Programs in a conventional language like C are written to access a database by “calls” to SQL statements.

# Host Languages

- ◆ Any conventional language can be a *host language*, that is, a language in which SQL calls are embedded.
- ◆ The use of a host/SQL combination allows us to do anything computable, yet still get the very-high-level SQL interface to the database.

# Connecting SQL to the Host Language

1. *Embedded SQL* is a standard for combining SQL with seven languages.
2. *CLI (Call-Level Interface)* is a different approach to connecting C to an SQL database.
3. *JDBC (Java Database Connectivity)* is a way to connect Java with an SQL database.

# Embedded SQL

- ◆ **Key idea:** Use a preprocessor to turn SQL statements into procedure calls that fit with the host-language code surrounding.
- ◆ All embedded SQL statements begin with EXEC SQL, so the preprocessor can find them easily.

# Shared Variables

- ◆ To connect SQL and the host-language program, the two parts must share some variables.
- ◆ Declarations of shared variables are bracketed by:

 `EXEC SQL BEGIN DECLARE SECTION;`

`<host-language declarations>`

 `EXEC SQL END DECLARE SECTION;`

# Use of Shared Variables

- ◆ In SQL, the shared variables must be preceded by a colon.
  - ◆ They may be used as constants provided by the host-language program.
  - ◆ They may get values from SQL statements and pass those values to the host-language program.
- ◆ In the host language, shared variables behave like any other variable.

# Example: Looking Up Prices

- ◆ We'll use C with embedded SQL to sketch the important parts of a function that obtains a beer and a bar, and looks up the price of that beer at that bar.
- ◆ Assumes database has our usual `Sells(bar, beer, price)` relation.



# Example: C Plus SQL

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
char theBar[21], theBeer[21];
```

```
float thePrice;
```

```
EXEC SQL END DECLARE SECTION;
```

```
/* obtain values for theBar and theBeer */
```

```
EXEC SQL SELECT price INTO :thePrice
```

```
FROM Sells
```

```
WHERE bar = :theBar AND beer = :theBeer;
```

```
/* do something with thePrice */
```

Note 21-char  
arrays needed  
for 20 chars +  
endmarker

SELECT-INTO  
just like PSM

# Embedded Queries

- ◆ Embedded SQL has the same limitations as PSM regarding queries:
  - ◆ You may use SELECT-INTO for a query guaranteed to produce a single tuple.
  - ◆ Otherwise, you have to use a cursor.
    - Small syntactic differences between PSM and Embedded SQL cursors, but the key ideas are identical.

# Cursor Statements

- ◆ Declare a cursor  $c$  with:

```
EXEC SQL DECLARE  $c$  CURSOR FOR <query>;
```

- ◆ Open and close cursor  $c$  with:

```
EXEC SQL OPEN CURSOR  $c$ ;
```

```
EXEC SQL CLOSE CURSOR  $c$ ;
```

- ◆ Fetch from  $c$  by:

```
EXEC SQL FETCH  $c$  INTO <variable(s)>;
```

- ◆ Macro NOT FOUND is true if and only if the FETCH fails to find a tuple.


# Example --- (1)

- ◆ Let's write C + SQL to print Joe's menu  
--- the list of beer-price pairs that we find in `Sells(bar, beer, price)` with `bar = Joe's Bar`.
- ◆ A cursor will visit each Sells tuple that has `bar = Joe's Bar`.

# Example --- (2: Declarations)

```
EXEC SQL BEGIN DECLARE SECTION;  
    char theBeer[21]; float thePrice;  
EXEC SQL END DECLARE SECTION;
```

```
EXEC SQL DECLARE c CURSOR FOR  
    SELECT beer, price FROM Sells  
    WHERE bar = 'Joe''s Bar';
```



The cursor declaration goes  
outside the declare-section

# Example --- (3: Executable)

```
EXEC SQL OPEN CURSOR c;
```

```
while(1){
```

```
    EXEC SQL FETCH c
```

```
        INTO :theBeer, :thePrice;
```

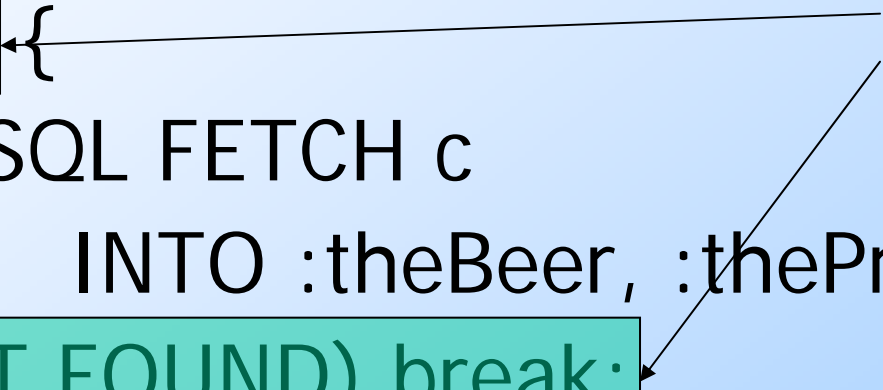
```
    if (NOT FOUND) break;
```

```
    /* format and print theBeer and thePrice */
```

```
}
```

```
EXEC SQL CLOSE CURSOR c;
```

The C style  
of breaking  
loops



# Need for Dynamic SQL

- ◆ Most applications use specific queries and modification statements to interact with the database.
  - ◆ The DBMS compiles EXEC SQL ... statements into specific procedure calls and produces an ordinary host-language program that uses a library.
- ◆ What about sqlplus, which doesn't know what it needs to do until it runs?

# Dynamic SQL

- ◆ Preparing a query:

```
EXEC SQL PREPARE <query-name>  
        FROM <text of the query>;
```

- ◆ Executing a query:

```
EXEC SQL EXECUTE <query-name>;
```

- ◆ "Prepare" = optimize query.
- ◆ Prepare once, execute many times.



# Example: A Generic Interface

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
    char query[MAX_LENGTH];
```

```
EXEC SQL END DECLARE SECTION;
```

```
while(1) {
```

```
    /* issue SQL> prompt */
```

```
    /* read user's query into array query */
```

```
    EXEC SQL PREPARE q FROM :query;
```

```
    EXEC SQL EXECUTE q;
```

```
}
```

q is an SQL variable  
representing the optimized  
form of whatever statement  
is typed into :query

# Execute-Immediate

- ◆ If we are only going to execute the query once, we can combine the PREPARE and EXECUTE steps into one.

- ◆ Use:

```
EXEC SQL EXECUTE IMMEDIATE <text>;
```

# Example: Generic Interface Again

```
EXEC SQL BEGIN DECLARE SECTION;  
    char query[MAX_LENGTH];  
EXEC SQL END DECLARE SECTION;  
while(1) {  
    /* issue SQL> prompt */  
    /* read user's query into array  
    query */  
    EXEC SQL EXECUTE IMMEDIATE :query;  
}
```

# SQL/CLI

- ◆ Instead of using a preprocessor, we can use a library of functions and call them as part of an ordinary C program.
  - ◆ The library for C is called SQL/CLI = "Call-Level Interface."
  - ◆ Embedded SQL's preprocessor will translate the EXEC SQL ... statements into CLI or similar calls, anyway.

# Data Structures

- ◆ C connects to the database by structs of the following types:
  1. *Environments* : represent the DBMS installation.
  2. *Connections* : logins to the database.
  3. *Statements* : SQL statements to be passed to a connection.
  4. *Descriptions* : records about tuples from a query or parameters of a statement.

# Environments, Connections, and Statements

- ◆ Function `SQLAllocHandle(T,I,O)` is used to create these structs, which are called environment, connection, and statement *handles*.
  - ◆  $T$  = type, e.g., `SQL_HANDLE_STMT`.
  - ◆  $I$  = input handle = struct at next higher level (statement < connection < environment).
  - ◆  $O$  = (address of) output handle.

# Example: SQLAllocHandle

```
SQLAllocHandle( SQL_HANDLE_STMT,  
               myCon, &myStat );
```

- ◆ myCon is a previously created connection handle.
- ◆ myStat is the name of the statement handle that will be created.

# Preparing and Executing

- ◆ **SQLPrepare( $H, S, L$ )** causes the string  $S$ , of length  $L$ , to be interpreted as an SQL statement and optimized; the executable statement is placed in statement handle  $H$ .
- ◆ **SQLExecute( $H$ )** causes the SQL statement represented by statement handle  $H$  to be executed.



# Example: Prepare and Execute

```
SQLPrepare(myStat, "SELECT beer, price  
FROM Sells WHERE bar = 'Joe''s Bar' ",  
SQL_NTS);  
SQLExecute(myStat);
```

This constant says the second argument is a "null-terminated string"; i.e., figure out the length by counting characters.

# Dynamic Execution

- ◆ If we will execute a statement  $S$  only once, we can combine PREPARE and EXECUTE with:

`SQLExecuteDirect(H,S,L);`

- ◆ As before,  $H$  is a statement handle and  $L$  is the length of string  $S$ .

# Fetching Tuples

- ◆ When the SQL statement executed is a query, we need to fetch the tuples of the result.
  - ◆ That is, a cursor is implied by the fact we executed a query, and need not be declared.
- ◆ **SQLFetch(H)** gets the next tuple from the result of the statement with handle *H*.

# Accessing Query Results

- ◆ When we fetch a tuple, we need to put the components somewhere.
- ◆ Thus, each component is bound to a variable by the function **SQLBindCol**.
  - ◆ This function has 6 arguments, of which we shall show only 1, 2, and 4:
    - 1 = handle of the query statement.
    - 2 = column number.
    - 4 = address of the variable.

# Example: Binding

- ◆ Suppose we have just done `SQLExecute(myStat)`, where `myStat` is the handle for query

```
SELECT beer, price FROM Sells  
WHERE bar = 'Joe''s Bar'
```

- ◆ Bind the result to `theBeer` and `thePrice`:  
`SQLBindCol(myStat, 1, , &theBeer, , );`  
`SQLBindCol(myStat, 2, , &thePrice, , );`

# Example: Fetching

- ◆ Now, we can fetch all the tuples of the answer by:

```
while ( SQLFetch(myStat) != SQL_NO_DATA)
{
    /* do something with theBeer and
       thePrice */
}
```

CLI macro representing  
SQLSTATE = 02000 = "failed  
to find a tuple."

# JDBC

- ◆ Java Database Connectivity (JDBC) is a library similar to SQL/CLI, but with Java as the host language.
- ◆ JDBC/CLI differences are often related to the object-oriented style of Java, but there are other differences.

# Environments, Connections, and Statements

- ◆ The same progression from environments to connections to statements that we saw in CLI appears in JDBC.
- ◆ A *connection object* is obtained from the environment in a somewhat implementation-dependent way.
- ◆ We'll start by assuming we have `myCon`, a connection object.



# Statements

- ◆ JDBC provides two classes:
  1. *Statement* = an object that can accept a string that is an SQL statement and can execute such a string.
  2. *PreparedStatement* = an object that has an associated SQL statement ready to execute.

# Creating Statements

- ◆ The Connection class has methods to create Statements and PreparedStatement.

```
Statement stat1 = myCon.createStatement();
```

```
PreparedStatement stat2 =
```

```
myCon.createStatement(
```

```
"SELECT beer, price FROM Sells "
```

```
"WHERE bar = 'Joe's Bar' "
```

```
);
```

Java trick: +  
concatenates  
strings.

+

`createStatement` with no argument returns a Statement; with one argument it returns a PreparedStatement.

# Executing SQL Statements

- ◆ JDBC distinguishes queries from modifications, which it calls “updates.”
- ◆ Statement and PreparedStatement each have methods `executeQuery` and `executeUpdate`.
  - ◆ For Statements, these methods have one argument: the query or modification to be executed.
  - ◆ For PreparedStatements: no argument.

# Example: Update

- ◆ stat1 is a Statement.

- ◆ We can use it to insert a tuple as:

```
stat1.executeUpdate(  
    "INSERT INTO Sells " +  
    "VALUES( 'Brass Rail', 'Bud', 3.00 )"  
);
```

# Example: Query

- ◆ stat2 is a PreparedStatement holding the query "SELECT beer, price FROM Sells WHERE bar = 'Joe''s Bar' ".
- ◆ `executeQuery` returns an object of class ResultSet --- we'll examine it later.
- ◆ The query:

```
ResultSet Menu = stat2.executeQuery();
```

# Accessing the ResultSet

- ◆ An object of type `ResultSet` is something like a cursor.
- ◆ Method `Next()` advances the “cursor” to the next tuple.
  - ◆ The first time `Next()` is applied, it gets the first tuple.
  - ◆ If there are no more tuples, `Next()` returns the value `FALSE`.

# Accessing Components of Tuples

- ◆ When a ResultSet is referring to a tuple, we can get the components of that tuple by applying certain methods to the ResultSet.
- ◆ Method `getX(i)`, where  $X$  is some type, and  $i$  is the component number, returns the value of that component.
  - ◆ The value must have type  $X$ .

# Example: Accessing Components

- ◆ Menu is the ResultSet for the query "SELECT beer, price FROM Sells WHERE bar = 'Joe's Bar'".
- ◆ Access the beer and price from each tuple by:

```
while ( Menu.Next() ) {  
    theBeer = Menu.getString(1);  
    thePrice = Menu.getFloat(2);  
    /* do something with theBeer and  
       thePrice */  
}
```