

CS 245: Database System Principles

Notes 11: View Serializability

Hector Garcia-Molina

CS 245

Notes 11

1

View Serializability

Conflict equivalent

View equivalent

Conflict serializable

View serializable

CS 245

Notes 11

2

Motivating example

Schedule Q

T ₁	T ₂	T ₃
Read(A)		
	Write(A)	
Write(A)		
		Write(A)


CS 245

Notes 11

3

Same as

$Q = r_1(A) \ w_2(A) \ w_1(A) \ w_3(A)$



►► Not conflict serializable!

CS 245

Notes 11

4

But now compare Q to S_s, a serial schedule:

Q	T ₁	T ₂	T ₃
	Read(A)		
	Write(A)	Write(A)	
			Write(A)

S _s	T ₁	T ₂	T ₃
	Read(A)		
	Write(A)	Write(A)	
			Write(A)

CS 245

Notes 11

5

- T₁ reads same thing in Q, S_s
- T₂, T₃ read something (nothing?)
- After Q or S_s, DB is left in same state

► So what is wrong with Q?

CS 245

Notes 11

6

Definition Schedules S_1, S_2 are View Equivalent if:

- (1) If in S_1 : $w_j(A) \Rightarrow r_i(A)$
then in S_2 : $w_j(A) \Rightarrow r_i(A)$
- (2) If in S_1 : $r_i(A)$ reads initial DB value,
then in S_2 : $r_i(A)$ also reads initial DB value
- (3) If in S_1 : T_i does last write on A ,
then in S_2 : T_i also does last write on A

\Rightarrow means "reads value produced"

CS 245

Notes 11

7

Definition

Schedule S_1 is View Serializable if it is view equivalent to some serial schedule

CS 245

Notes 11

8

View Serializable $\xleftarrow{\quad} ? \xrightarrow{\quad}$ Conflict Serializable

- View Serializable $\not\Rightarrow$ Conflict Serializable
e.g., See Schedule Q
- Conflict Serializable $\stackrel{?}{\Rightarrow}$ View Serializable

CS 245

Notes 11

9

Lemma

Conflict Serializable \Rightarrow View Serializable

Proof:

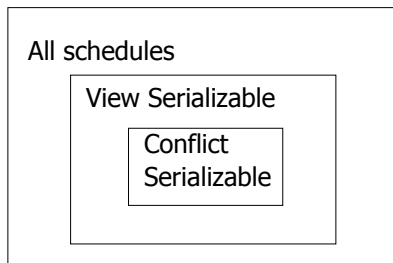
Swapping non-conflicting actions does not change what transactions read nor final DB state

CS 245

Notes 11

10

Venn Diagram



CS 245

Notes 11

11

Note: All view serializable schedules that are not conflict serializable, involve useless write

$S = W_2(A) \dots W_3(A) \dots$
no reads

FALSE: Counterexample (Sorav Bansal):
 $w_3(Y) r_2(Y) w_1(X) r_2(X) w_3(X) r_4(X) w_5(X)$

CS 245

Notes 11

12

How do we test for view-serializability?

⇒ P(S) not good enough...
(see schedule Q)

CS 245

Notes 11

13

- One problem: some swaps involving conflicting actions are OK... e.g.:

$S = \dots w_2(A) \dots r_1(A) \dots w_3(A) \dots w_4(A)$

Diagram: A dashed arrow points from $w_3(A)$ to $w_2(A)$ with the text "this action can move if this write exists".

CS 245

Notes 11

14

- Another problem: useless writes

$S = \dots w_2(A) \dots w_1(A) \dots$

Diagram: A double-headed arrow between $w_2(A)$ and $w_1(A)$ with the text "no A reads" below it.

CS 245

Notes 11

15

To check if S is View Serializable

- (1) Add final transaction T_f that reads all DB
(eliminates condition 3 of V-S definition)

E.g.: $S = \dots w_1(A) \dots w_2(A) \dots r_f(A)$

Diagram: A curved arrow with a question mark points from $w_2(A)$ to $r_f(A)$. Below $w_2(A)$ is the label "Last A write" with an upward arrow. Below $r_f(A)$ is the label "add" with an upward arrow.

CS 245

Notes 11

16

- (2) Add initial transaction T_b that writes all DB
(eliminates condition 2 of V-S definition)

E.g.: $S = w_b(A) \dots r_1(A) \dots w_2(A) \dots$

Diagram: An upward arrow points to $w_b(A)$ with the label "add". A curved arrow with a question mark points from $w_2(A)$ back to $w_b(A)$.

CS 245

Notes 11

17

- (3) Create labeled precedence graph of S:
(3a) If $w_i(A) \Rightarrow r_j(A)$ in S, add $T_i \xrightarrow{A} T_j$

CS 245

Notes 11

18

- (3b) For each $w_i(A) \Rightarrow r_j(A)$ do
 consider each $w_k(A)$: $[T_k \neq T_b]$
- If $T_i \neq T_b \wedge T_j \neq T_f$ then insert

$$\begin{cases} T_k \xrightarrow{p} T_i \\ T_j \xrightarrow{p} T_k \end{cases} \quad \text{some new } p$$
 - If $T_i = T_b \wedge T_j \neq T_f$ then insert

$$T_j \xrightarrow{0} T_k$$
 - If $T_i \neq T_b \wedge T_j = T_f$ then insert

$$T_k \xrightarrow{0} T_i$$

CS 245

Notes 11

19

- (4) Check if $LP(S)$ is "acyclic" (if so, S is V-S)
 - For each pair of "p" arcs ($p \neq 0$),
 choose one

CS 245


Notes 11

20

Example: check if Q is V-S:

$Q = r_1(A) w_2(A) w_1(A) w_3(A)$

$Q' = w_b(A) \Rightarrow r_1(A) w_2(A) w_1(A) w_3(A) \Rightarrow r_f(A)$



rule 3(a)
 rule 3(b)
 rule 3(b)

LP(S) acyclic!!
 S is V-S


CS 245

Notes 11

21

Another example:

$Z = w_b(A) \Rightarrow r_1(A) w_2(A) \Rightarrow r_3(A) w_1(A) w_3(A) \Rightarrow r_f(A)$



CS 245

Notes 11

22

$S_s = w_b(A) \underbrace{r_1(A) w_1(A)}_{T_1} \underbrace{w_2(A) r_3(A)}_{T_2} \underbrace{w_3(A) r_f(A)}_{T_3}$

$Z + S_s$ indeed do same thing

- Checking view serializability is expensive
- Still, V-S useful in some cases...

CS 245

Notes 11

23

CS 245

Notes 11

24

Example on useless transactions:

$S = w_1(A) \ r_2(A) \ w_2(B) \ r_1(B) \ w_3(A) \ w_3(B)$


CS 245

Notes 11

25

$S' =$

$T_b \ w_1(A) \Rightarrow r_2(A) \ w_2(B) \Rightarrow r_1(B) \ w_3(A) \ w_3(B) \Rightarrow T_f$



CS 245

Notes 11

26

- If we only care about final state ➡ remove T_1, T_2 ; i.e., remove useless transactions
- If we care what T_1, T_2 read (view equivalence), then do not remove useless transactions

CS 245

Notes 11

27

- If all transactions read what they write, (I.e., $T_j = \dots R_j(A) \dots W_j(A) \dots$) then view serializability = conf. serializability

[Another way of saying: blind writes appear in any view-serializable schedule that is not conflict serializable]

CS 245

Notes 11

28

Proof(?): say S_1 is view-ser. and no blind writes. S_1 V-equiv to S_s , serial schedule.

- (1) Goal: Show that $T_1 \rightarrow T_2$ in $P(S_1) \Rightarrow T_1 <_{ss} T_2$
 - (2) Assume $T_1 \rightarrow T_2$
 - if $S_1 = \dots w_1(A) \dots r_2(A) \dots$
(direct read) clearly $T_1 <_{ss} T_2$
 - if $S_1 = \dots w_1(A) \dots r_3(A) \ w_3(A) \dots r_2(A) \dots$
also $T_1 <_{ss} T_2$
 - if $S_1 = \dots r_1(A) \ r_3(A) \dots w_1(A) \dots w_3(A) \dots r_2(A) \dots$
not possible: T_1, T_3 not serializable
- Other cases similar...

CS 245

Notes 11

29

Implications:

If no blind writes, view-ser \iff conf-ser

$P(S)$ acyclic \Rightarrow all transactions read the same as in a serial schedule

CS 245

Notes 11

30